

Strategic Architecture and Computational Framework for the Generative Eyewear Design Studio

1. Strategic Context and Analytical Directives

The transition from conventional eyewear manufacturing to a digitally native, computationally driven design ecosystem requires the synthesis of generative artificial intelligence, parametric computer-aided design (CAD), and advanced computer-aided manufacturing (CAM). Within the broader framework of developing an intelligent eyewear design platform, Phase 3 (Interactive Studio: Sketch-to-Color) and Phase 4 (Advanced Generation & Automation) represent the critical junctures where raw trend telemetry and aesthetic concepts are translated into physical, manufacturable products¹.

Phase 3 demands an interactive graphical environment where designers can upload monochromatic 2D sketches, delineate specific lamination zones, and apply photorealistic material textures derived from actual supplier catalogs, such as those from Mazzucchelli 1849¹. Phase 4 requires the deployment of an autonomous generative engine capable of synthesizing market trend data and brand-specific aesthetic codes (Brand DNA) to propose novel frame variations that are ready for prototyping and industrial production¹.

This report provides an exhaustive analysis of the technological prerequisites for these phases. It addresses three primary directives: the identification of analogous projects and market precedents, the definitive necessity of Large Language Models (LLMs) and Image Models, and the comprehensive software stack required to bridge generative AI with high-precision CNC (Computer Numerical Control) milling.

2. Market Precedents and Analogous Technological Projects

While a fully integrated, end-to-end generative AI studio specifically tailored for physically manufactured acetate eyewear is a highly novel proposition, various components of the required technology are actively being pioneered across adjacent sectors. Analyzing these precedents validates the technological feasibility of the proposed studio and highlights critical architectural patterns.

2.1. Industrial Sketch-to-Render Platforms

The foundational requirement of Phase 3—converting rough 2D linework into a photorealistic render while strictly preserving the underlying geometry—is currently being spearheaded by specialized industrial design platforms.

Vizcom stands out as a leading AI-powered rendering tool that accelerates product development from concept to prototype³. The platform enables industrial designers to visualize form, materials, and lighting instantaneously by utilizing advanced neural networks that respect

edges, proportions, and original linework, allowing users to apply region masks for specific material renderings⁴. Similarly, platforms like Floyo.ai utilize specialized "Lineart" and "Previz" workflows to convert rough sketches into polished, colorized concepts, providing designers with a range of unique model outputs suitable for establishing visual consistency⁶. These platforms demonstrate that pixel-perfect structural retention during the diffusion generation process is computationally viable and highly effective for industrial ideation.

2.2. Eyewear-Specific Virtual Try-On (VTO) and Augmented Reality

In the consumer-facing eyewear market, generative AI and computer vision are heavily utilized to create 3D assets from 2D inputs, a mechanism directly applicable to the rendering engine required for the proposed studio.

The platform Jeeliz utilizes generative artificial intelligence to automatically create detailed 3D models of glasses directly from 2D product page pictures, standardizing lighting to match the user's face for unmatched realism in Virtual Try-On (VTO) applications⁷. Makeover.so operates on a similar paradigm, utilizing computer vision and generative AI to scale frames to a user's exact facial proportions—measuring pupillary distance, cheekbone height, and jawline curves—to generate photorealistic previews⁸.

Furthermore, Kivicube utilizes generative AI to automate 3D content generation for WebAR, completely shifting the paradigm of traditional 3D modeling⁹. Kivicube's engines can turn simple text prompts or 2D images into AR-ready 3D assets featuring physically based rendering (PBR) textures that perfectly replicate the visual characteristics of real acetate frames and lenses⁹. This capability to accurately synthesize the translucency and reflectivity of acetate is a critical prerequisite for Phase 3¹.

2.3. Algorithmic Aesthetic Decoding and Brand DNA

Phase 4 requires the AI to understand and autonomously generate designs that align with a specific brand's stylistic signatures (Brand DNA)¹. This moves the system from a passive rendering tool to an active creative participant.

A highly relevant precedent is the collaboration between British eyewear brand Black Eyewear and Google Cloud¹⁰. Facing the challenge of scaling the highly subjective, refined taste of its founder, the brand trained specialized models on curated datasets to see if AI could replicate elite creative judgment¹⁰. The project demonstrated that AI can move beyond being an "averaging machine" to actively preserve and extend elite creative expertise, a capability essential for maintaining the distinct aesthetic boundaries of brands like *Noe Noah*¹. Similarly, platforms like InAirSpace leverage machine learning to decode personal aesthetics, moving from pure biometric measurement to understanding qualitative style parameters like "bold," "minimalist," or "vintage"¹¹.

Precedent Platform	Core Technology Utilized	Application to the Generative Studio (Phases 3 & 4)

Vizcom / Floyo.ai	ControlNet, Edge Preservation, Region Masking	Validates the Sketch-to-Lamination workflow; ensures 2D sketches maintain geometric accuracy during rendering ³ .
Kivicube	Generative 3D AR, PBR Texture Synthesis	Demonstrates the ability to synthesize physically accurate acetate and lens reflections from 2D/text inputs ⁹ .
Makeover.so / Jeeliz	Computer Vision, 2D-to-3D Extrapolation	Provides frameworks for extracting precise spatial data and standardizing lighting across generated frame models ⁷ .
Black Eyewear (Google Cloud)	Custom Model Fine-Tuning, Aesthetic Decoding	Proves the viability of training generative models to adhere strictly to elite, proprietary Brand DNA ¹⁰ .

3. The Necessity and Architecture of Image Models

To address the core strategic query: **The deployment of an advanced, composite Image Model architecture is absolutely and strictly necessary.**

Standard, off-the-shelf text-to-image models (e.g., standard Midjourney or DALL-E) are fundamentally incompatible with industrial eyewear design. They operate on a stochastic basis, frequently "hallucinating" geometries, violating mechanical constraints, and failing to adhere to the rigid dimensional tolerances required for physical manufacturing¹. Eyewear design is an exercise in millimeters; a slight deviation in the bridge or tenon renders a frame unwearable or un-millable¹.

Therefore, the studio must utilize open-weights latent diffusion models (such as Stable Diffusion XL or Flux) heavily augmented with deterministic neural conditioning adapters.

3.1. Structural Conditioning: ControlNet and SAM

To execute Phase 3, the generative engine must ingest a designer's 2D vector or bitmap sketch and generate a render without altering the silhouette¹. This requires **ControlNet**, a neural network architecture designed to add spatial conditioning controls to large, pretrained text-to-image diffusion models¹³. By freezing the original weights of the diffusion model and creating a trainable copy via zero convolutions, ControlNet injects structural priors without

destroying the model's base generative capabilities¹⁴.

Specific ControlNet models are required for the eyewear pipeline:

- **Lineart and Canny Models:** These models extract high-frequency geometric contours from the input sketch and force the latent noise generation to strictly follow these boundaries¹⁵. This guarantees that the shape of the lenses and the curvature of the frame face remain identical to the designer's intent.
- **Depth Estimation Models:** Eyewear frames are three-dimensional objects. Depth models analyze spatial relationships, differentiating the foreground (frame face) from the midground (temples) and background, which is essential for simulating the thickness of acetate laminations¹⁵.

Furthermore, to facilitate the manual delineation of lamination zones (e.g., specifying a different acetate for the lower rim versus the upper browline), the architecture must incorporate the **Segment Anything Model (SAM)**¹. SAM provides zero-shot image segmentation, allowing the designer to interactively select regions of the sketch to generate precise masks¹⁷. These masks constrain the subsequent material diffusion strictly to the designated areas, enabling complex bi-layer or tri-layer acetate renderings¹.

3.2. Material Injection: IP-Adapter and Texture Synthesis

The studio's value proposition relies on rendering frames using actual, procurable materials from suppliers like Mazzucchelli 1849, rather than AI-hallucinated colors¹. Textual prompts are fundamentally inadequate for this task; describing a "translucent Havana tortoise shell with fine amber marbling" will yield inconsistent and inaccurate results.

The solution is the **Image Prompt Adapter (IP-Adapter)**. This lightweight adapter introduces a decoupled cross-attention mechanism that works in parallel with the standard text cross-attention layers in the diffusion model²⁰.

- **Mechanism:** When a designer selects a specific acetate swatch from the digitized supplier catalog, the IP-Adapter uses an image encoder (like CLIP) to project the visual embeddings of that swatch directly into the denoising process²⁰.
- **Application:** The diffusion model effectively "paints" the exact visual texture, color distribution, and structural pattern of the physical acetate onto the geometry locked by the ControlNet²³.

3.3. Physically Based Rendering (PBR) Latent Generation

Acetate interacts dynamically with light through subsurface scattering, specular reflection, and varying degrees of opacity. Standard RGB image generation cannot capture these properties accurately. To achieve true industrial fidelity, the Image Model stack must be capable of generating **Physically Based Rendering (PBR)** materials²⁵.

Recent advancements in texture synthesis, such as TexGaussian, DreamMat, and MatLat, address this by forcing diffusion models to generate separated material maps—specifically Albedo (base color), Roughness, and Metallic/Specular channels—rather than baked-in RGB lighting²⁵. By utilizing geometry- and light-aware diffusion models, the studio can synthesize the high-gloss polish of premium cellulose acetate or the matte finish of 3D-printed polyamide without unwanted baked-in shadows, ensuring the render accurately reflects how the physical

product will interact with real-world environments²⁵.

3.4. Brand DNA Encoding: Low-Rank Adaptation (LoRA)

Phase 4 mandates that the AI autonomously generates designs that fit within a brand's existing portfolio¹. This is achieved through **LoRA (Low-Rank Adaptation)**.

LoRA is a parameter-efficient fine-tuning technique that injects trainable rank decomposition matrices into the frozen layers of a pre-trained diffusion model²⁹. By compiling a dataset of a brand's historical models, technical drawings, and aesthetic signatures, a custom LoRA can be trained rapidly³¹. When applied during generation, the LoRA steers the model's output to inherently adopt the brand's preferred proportions, hinge styles, and geometric flow, ensuring the autonomous outputs are highly curated rather than generic¹⁴.

Model / Adapter	Architectural Function	Phase Application
Stable Diffusion XL / Flux	Base latent diffusion foundation for pixel generation.	Phases 3 & 4
ControlNet (Lineart/Depth)	Constrains pixel generation to exact sketch boundaries.	Phase 3
Segment Anything (SAM)	Generates precise spatial masks for lamination zones.	Phase 3
IP-Adapter	Bypasses text prompts to directly inject photographic material textures.	Phases 3 & 4
PBR Synthesis Models	Generates Albedo, Roughness, and Metallic maps for accurate lighting.	Phases 3 & 4
LoRA	Fine-tunes the base model to adhere to specific Brand DNA parameters.	Phase 4

4. The Necessity and Architecture of Large Language Models (LLMs)

While Image Models handle the spatial and textural synthesis, a **Large Language Model (LLM)**

is strictly necessary to serve as the cognitive orchestration engine, translating statistical market data into actionable creative directives. Image models cannot natively parse trend statistics, filter supplier databases, or reason about manufacturing constraints; the LLM bridges this gap.

4.1. Semantic Orchestration and Trend Translation

Phases 1 and 2 of the overarching project involve the continuous ingestion of market telemetry via a Trend Radar, tracking the popularity of specific shapes, materials, and colors across global markets¹. This data is structured and quantitative.

The LLM is required to analyze this telemetry and synthesize it into generative parameters. For example, if the Trend Radar detects a statistically significant increase in the adoption of geometric titanium-acetate hybrid frames in the European market, the LLM interprets this trend and formulates the specific semantic instructions required to drive the image generation pipeline in Phase 4¹.

4.2. Retrieval-Augmented Generation (RAG)

To ensure the LLM's outputs are firmly grounded in reality and physical supply chains, it must be integrated with a **Retrieval-Augmented Generation (RAG)** architecture³³. RAG connects the LLM to external, proprietary vector databases, feeding verified, domain-specific data into the model's context window prior to generation³⁵.

In the context of the generative studio, the RAG system will index:

1. **The Material Database:** Digitized supplier catalogs (e.g., Mazzucchelli acetate codes, thicknesses, and availability)¹.
2. **Trend Telemetry:** Real-time data from the Phase 1 Radar¹.
3. **Manufacturing Constraints:** Rules defining the minimum thickness required for lamination, the necessary mechanical clearances for tenon hinges, and CNC milling tolerances¹.

When the LLM is prompted to suggest a new design, it first retrieves the available in-stock acetates and cross-references them against the trending color palettes, ensuring it never proposes a material that cannot be physically procured or manufactured.

4.3. Automated Prompt Engineering (APE)

In Phase 4, the LLM acts autonomously through **Automated Prompt Engineering (APE)**³⁶. APE allows the LLM to construct and optimize highly complex, multimodal prompts for the diffusion backend³⁶.

The LLM algorithmically compiles a multi-layered directive: it selects a base shape influenced by the brand's LoRA, retrieves the exact image embedding of an in-stock acetate to feed into the IP-Adapter, defines the geometric boundaries for the ControlNet, and sets the negative prompts to avoid common visual artifacts¹⁵. This transforms the LLM into a fully automated art director, orchestrating the various image adapters to produce highly curated, trend-aligned capsule collections¹.

5. Software Stack: Interactive Frontend Canvas

The user interface for Phase 3 must provide a highly responsive, intuitive environment where

designers can upload sketches, draw bounding boxes, manipulate splines for lamination zones, and assign materials. This requires a sophisticated frontend graphics library.

5.1. WebGL vs. HTML5 Canvas

The architectural decision fundamentally rests between WebGL and HTML5 Canvas APIs³⁷.

- **WebGL:** A JavaScript API that brings hardware-accelerated, GPU-powered graphics to the browser. It excels in handling complex 3D models, large datasets, and real-time rendering of advanced lighting and shading³⁷.
- **Canvas:** A simpler, CPU-based API optimized for 2D graphics, basic interactivity, and rapid prototyping³⁷.

For the primary 2D interaction of Phase 3 (sketch uploading, vector drawing, and masking), **HTML5 Canvas** is the most pragmatic choice. While WebGL offers superior performance for 3D assemblies, the initial interaction involves 2D vectors and bitmaps, for which Canvas provides smoother native text rendering, easier DOM event handling, and seamless integration with specialized 2D libraries³⁷. However, if the studio expands to allow users to rotate the AI-generated frame in 3D space to inspect PBR reflections, a hybrid approach utilizing a WebGL wrapper like Three.js for the final output viewer will be necessary³⁹.

5.2. Canvas Libraries: Fabric.js vs. Konva.js

Within the Canvas ecosystem, the two dominant libraries for building interactive graphic editors are Fabric.js and Konva.js.

- **Konva.js:** Known for its high performance and declarative nature, Konva is excellent for applications requiring complex animations or deep node hierarchies³⁷. However, it requires developers to build many interactive shortcuts from scratch³⁹.
- **Fabric.js:** Fabric provides a highly robust object-oriented model specifically designed for interactive drawing³⁷. It features built-in support for image filters, free-hand drawing, complex vector manipulation, and critically, native serialization to and from SVG (Scalable Vector Graphics)³⁹.

Given that the core interaction in Phase 3 involves designers delineating lamination zones over a sketch, **Fabric.js** is the optimal choice. Its out-of-the-box support for interactive bounding boxes, bezier curve manipulation, and SVG serialization perfectly aligns with the requirement to translate designer inputs into precise masks for the SAM and ControlNet models³⁷.

6. Software Stack: Backend, CAD, and CAM

Integration

The most critical challenge of the entire project is ensuring "faisabilité technique" (technical feasibility)—bridging the gap between the AI-generated image and the physical CNC milling of cellulose acetate¹. An AI rendering is useless to an industrialist if it cannot be converted into a machine toolpath.

6.1. Backend Orchestration

The backend must handle the asynchronous, computationally heavy inference of the diffusion models and the LLM.

- **API Framework:** A Python-based framework, such as FastAPI, is standard due to its high performance and native compatibility with machine learning ecosystems like PyTorch.
- **Inference Middleware:** To manage the complex chaining of models (SDXL + ControlNet + IP-Adapter + LoRA), the backend should utilize a graph-based execution engine like **ComfyUI** or Hugging Face's **Diffusers** library⁴¹. These frameworks allow developers to programmatically define the generative pipeline via API calls, dynamically adjusting the weights of the IP-Adapter or the strength of the ControlNet based on the user's frontend input²⁰.

6.2. Parametric CAD Interoperability

In the eyewear industry, physical frames are engineered using parametric CAD software, with **Rhinoceros 3D (Rhino)** and its visual programming plugin **Grasshopper** being the undisputed industry standards⁴³. Grasshopper allows designers to build algorithmic definitions where dimensions are mathematically dependent, enabling rapid iterations of frame sizes, base curves, and lamination thicknesses⁴⁴.

To move from the web-based studio to Rhino, the system requires precise vector data exchange⁴⁶. The frontend Fabric.js canvas must export the designer's validated sketch and lamination boundaries as **SVG (Scalable Vector Graphics)** or **DXF (Drawing Exchange Format)** files⁴⁷. These formats preserve the mathematical curves (NURBS and polylines) rather than pixels, allowing them to be imported directly into Rhino⁴⁶.

Once in Rhino, these 2D vectors are parametrically extruded into 3D solid models or surfaces (B-rep), representing the precise thickness of the acetate sheets⁴⁶. The Grasshopper algorithms can automatically apply necessary manufacturing tolerances, such as the draft angles required for milling tools and the specific embossing cutouts for the temple core wires⁴⁴.

6.3. CAM and CNC Milling (G-Code Generation)

The final stage of the software stack involves translating the 3D CAD model into instructions for the manufacturing hardware.

- **Computer-Aided Manufacturing (CAM):** The Rhino models are processed through CAM software. The CAM software analyzes the geometry and calculates the optimal toolpaths for the milling cutters to remove excess material from the raw Mazzucchelli acetate blanks¹².
- **CNC Machining:** The CAM software outputs **G-Code**, the standardized programming language that controls CNC (Computer Numerical Control) machines⁵¹. The G-Code dictates the exact Cartesian coordinates, feed rates, and spindle speeds for the 3-axis or 5-axis CNC milling machines⁵¹.

High-speed CNC milling is the primary method for manufacturing premium acetate frames, allowing for complex 3D reliefs, gradient cuts, and precise lamination assemblies¹². By ensuring the frontend studio exports clean DXF/SVG data, the platform guarantees that the AI-assisted designs flow seamlessly into standard industrial CAM pipelines, preventing the "digital dead-end" common to many generative art platforms⁴⁶.

7. Execution Strategy and Implementation Roadmap

Synthesizing these technological components yields a clear, executable roadmap for deploying Phases 3 and 4, ensuring the outputs are both aesthetically superior and industrially viable.

7.1. Executing Phase 3: Interactive Sketch-to-Lamination Studio

1. **Material Ingestion:** Digitize physical supplier catalogs (e.g., Mazzucchelli 1849 PDFs)¹. Extract high-resolution swatches and convert them into image embeddings optimized for the IP-Adapter²⁰.
2. **Frontend Interaction:** The designer uploads a 2D DXF or bitmap sketch to the React/Fabric.js canvas¹. Using intuitive tools, the designer draws splines to delineate lamination zones.
3. **Semantic Masking:** The backend utilizes the Segment Anything Model (SAM) to perfectly map the designer's rough inputs to the exact geometric boundaries of the sketch, generating precise alpha masks for each frame component¹.
4. **Generative Inference:** The designer selects an acetate material from the database. The API triggers the ComfyUI pipeline. ControlNet Lineart locks the geometry, while the IP-Adapter injects the specific acetate texture solely within the SAM-defined mask¹⁵. The model outputs a photorealistic PBR render showcasing the material's true specular highlights and thickness¹.
5. **Parametric Export:** Upon approval, the Fabric.js canvas serializes the vector boundaries and exports them as a DXF file, ready for immediate import into Rhino/Grasshopper for 3D extrusion and CAM processing⁴⁷.

7.2. Executing Phase 4: Autonomous Generation and Automation

1. **Telemetry Ingestion:** The LLM continuously ingests statistical data from the Phase 1/2 Trend Radar, identifying emerging shifts in color preferences, morphological shapes, and lamination techniques within the target markets¹.
2. **RAG-Assisted Ideation:** The LLM queries the vector database via Retrieval-Augmented Generation to identify available, in-stock materials that align with the identified trends, ensuring supply chain viability³³.
3. **Automated Prompt Construction:** Utilizing Automated Prompt Engineering, the LLM constructs the complex JSON payload for the diffusion engine³⁶. It selects a base geometry profile, retrieves the appropriate acetate image embeddings for the IP-Adapter, and formulates the negative prompts.
4. **Brand DNA Application:** The LLM activates the brand-specific LoRA model, ensuring the generative output conforms to the proprietary stylistic signatures, hinge geometries, and aesthetic codes of the target brand¹.
5. **Industrial Output:** The system autonomously renders a highly curated capsule collection of variants. These are presented to the creative director with attached trend justification, fully prepped for DXF export and physical CNC prototyping, completing the cycle from raw market data to physical product¹.

By strictly adhering to this architectural framework, the generative design studio bridges the

computational capabilities of modern artificial intelligence with the rigid physical constraints of industrial eyewear manufacturing.

Works cited

1. cadrage_projet_radar_tendances_studio_ia_lunetterie.md
2. Materials For Design | PDF | Wood | Oak - Scribd,
<https://www.scribd.com/document/352592415/Materials-for-Design>
3. Vizcom for Industrial Design | Sketch to Render in Seconds,
<https://vizcom.com/solutions/industrial-design>
4. Vizcom Features | AI Rendering, Collaboration & Design Tools,
<https://vizcom.com/product>
5. Vizcom AI: Sketch to Render and 3D Models for Designers - AiTing AI,
<https://aitools.aiting.com/ai/vizcom>
6. floyoofficial, <https://www.floyo.ai/creators/floyoofficial>
7. Eyewear Virtual Try-on for the Web, <https://jeeliz.com/>
8. Virtual Glasses Try On: Find Your Frame Shape Free - Makeover,
<https://www.makeover.so/blog/guide/virtual-glasses-try-on>
9. Kivicube AI AR: Where Generative AI Meets Augmented Reality,
<https://www.kivicube.com/post/gen-ai-for-augmented-reality-how-to-build-immersive-experiences-without-coding/>
10. Black Eyewear case study | Google Cloud,
<https://cloud.google.com/customers/black-eyewear>
11. AI Choose My Glasses: The Ultimate Guide to Algorithmic Eyewear,
<https://inairspace.com/blogs/learn-with-inair/ai-choose-my-glasses-the-ultimate-guide-to-algorithmic-eyewear-selection>
12. CAD/CAM Principles and Applications Textbook - Studylib,
<https://studylib.net/doc/27048695/p.-n.-rao---cad-cam--principles-and-applications--3rd-edn...>
13. RefAdGen: High-Fidelity Advertising Image Generation,
<https://ojs.aaai.org/index.php/AAAI/article/download/37307/41269>
14. GAPS: Generative Active Pseudo-view Selection for Sparse ... - arXiv,
<https://arxiv.org/html/2609.23436>
15. A Cross-Stage Human-GenAI Collaborative Workflow for Landscape,
<https://www.mdpi.com/2075-5309/16/19/3819>
16. Paper Digest: Recent Papers on Style Transfer,
<https://www.paperdigest.org/2020/06/recent-papers-on-style-transfer/>
17. Daily Papers - Hugging Face, <https://huggingface.co/papers?q=reference-based>
18. Help! Is there a way to get real object exactly into a generated image?,
https://www.reddit.com/r/StableDiffusion/comments/18y87ud/help_is_there_a_way_to_get_real_object_exactly/
19. nnanwube/controlnet-brand-fidelity-benchmark - Hugging Face,
<https://huggingface.co/datasets/nnanwube/controlnet-brand-fidelity-benchmark>
20. IPAdapter Tutorial - Comflowy,
<https://www.comflowy.com/blog/IPAdapter-Tutorial>

21. Efficient and Controllable Image Generation on the Edge: A Survey,
<https://www.mdpi.com/2079-9292/15/4/828>
22. Text Compatible Image Prompt Adapter for Text-to-Image Diffusion,
https://www.researchgate.net/publication/373116447_IP-Adapter_Text_Compatible_Image_Prompt_Adapter_for_Text-to-Image_Diffusion_Models
23. DecorAI - AI Interior Design | Room Redesign App, <https://www.mydecor.app/>
24. ICAS: IP Adapter and ControlNet-based Attention Structure for Multi,
https://www.researchgate.net/publication/390959314_ICAS_IP_Adapter_and_ControlNet-based_Attention_Structure_for_Multi-Subject_Style_Transfer_Optimization
25. TexGaussian: Generating High-quality PBR Material via Octree,
https://openaccess.thecvf.com/content/CVPR2025/papers/Xiong_TexGaussian_Generating_High-quality_PBR_Material_via_Octree-based_3D_Gaussian_Splatting_CVPR_2025_paper.pdf
26. (PDF) MatLat: Material Latent Space for PBR Texture Generation,
https://www.researchgate.net/publication/398937058_MatLat_Material_Latent_Space_for_PBR_Texture_Generation
27. DreamMat: High-quality PBR Material Generation with Geometry,
<https://zzzyuqing.github.io/dreammat.github.io/>
28. Physically Consistent PBR Material Estimation at Interactive Rates,
https://openaccess.thecvf.com/content/ICCV2025/papers/Hong_SuperMat_Physically_Consistent_PBR_Material_Estimation_at_Interactive_Rates_ICCV_2025_paper.pdf
29. Lora Finetuning Guide - a Hugging Face Space by Bar-Fin,
<https://huggingface.co/spaces/Bar-Fin/lora-finetuning-guide>
30. Mastering AI: Model Training, Fine-Tuning, and LoRA Explained,
<https://www.blackcubelabs.com/blog/mastering-ai-model-training-fine-tuning-and-lora-explained>
31. Corporate AI Image Generator - Stellans,
<https://stellans.io/corporate-ai-image-generator/>
32. Harnessing the Power of Foundation Models for your Business with,
<https://www.ml6.eu/en/blog/harnessing-the-power-of-foundation-models-for-your-business-with-lora>
33. What is RAG? - Retrieval-Augmented Generation AI Explained - AWS,
<https://aws.amazon.com/what-is/retrieval-augmented-generation/>
34. Emerging Patterns in Building GenAI Products - Martin Fowler,
<https://martinfowler.com/articles/gen-ai-patterns/>
35. RAG vs fine-tuning vs. prompt engineering - IBM,
<https://www.ibm.com/think/topics/rag-vs-fine-tuning-vs-prompt-engineering>
36. What Is Automated Prompt Engineering? Guide & Uses - Truefoundry,
<https://www.truefoundry.com/glossary/automated-prompt-engineering>
37. WebGL vs Canvas: Best Choice for Browser-Based CAD Tools,
<https://altersquare.io/blog/webgl-vs-canvas-best-choice-for-browser-based-cad-tools>
38. WebGL vs Canvas: Best Choice for Browser-Based CAD Tools,
<https://altersquare.medium.com/webgl-vs-canvas-best-choice-for-browser-bas>

[ed-cad-tools-231097daf063](#)

39. [AskJS] What is the best canvas library to make an app like figma or,
https://www.reddit.com/r/javascript/comments/pdib99/askjs_what_is_the_best_canvas_library_to_make_an/
40. Best canvas library for React? : r/reactjs - Reddit,
https://www.reddit.com/r/reactjs/comments/qoq2m2/best_canvas_library_for_react/
41. A guide for beginners on a controlnet workflow (Nothing new for,
https://www.reddit.com/r/StableDiffusion/comments/1167j0a/a_guide_for_beginners_on_a_controlnet_workflow/
42. GitHub - tencent-ailab/IP-Adapter: The image prompt adapter is,
<https://github.com/tencent-ailab/IP-Adapter>
43. CNC Machining Line – Precision for Metal & Acetate Frames,
<https://eyewearbeyond.com/killing-page/cnc-machining-line-precision-for-metal-acetate-frames/>
44. Automating Product Design and Fabrication Within the Furniture,
<https://www.mdpi.com/2411-9660/9/5/116>
45. 1 2 This doctoral thesis offers a profound exploration of ... - IRIS UniGe,
https://unige.iris.cineca.it/retrieve/5118bdf0-e6ba-48c4-8ab4-b01e65f9c3a0/phd_unige_4958989.pdf
46. GitHub - naranyala/awesome-3d-modeling-for-cad: A curated list of,
<https://github.com/naranyala/awesome-3d-modeling-for-cad>
47. 3D Modeling in CAD: Complete Guide 2026 | SimuTecra,
<https://simutecra.com/blogs/3d-modeling-in-cad-complete-guide>
48. SolidVents | Parametric Pattern Generator — SVG, DXF & STEP,
<https://solidvents.com/en/>
49. China Cheap Full-Rim Round Frames Manufacturers Factory,
<https://www.songraineyewear.com/optical-frames/full-rim-eyewear/full-rim-round-frames.html>
50. Design and development of 3D printed fins integrated into an,
<https://ro.uow.edu.au/ndownloader/files/50392377/1>
51. 3D Printed Track and Turnouts - Templot Club,
https://85a.uk/templot/archive/topics/topic_2734.php
52. CNC Programming Handbook A Comprehensive Guide To Practical,
<https://www.scribd.com/document/1021770006/CNC-Programming-Handbook-A-Comprehensive-Guide-to-Practical-CNC-Programming-3rd-Edition-Edition-Smid-official-digital-release>
53. CAD/CAM Course Material Overview | PDF | Computer Aided Design,
<https://www.scribd.com/document/461232762/COMPUTER-AIDED-DESIGN-COMPUTER-AIDED-MANUFACTURING-pdf>